

Lecture 11:

Linear Filters I:

System Theory

Contents

1. Motivation
2. Linear System Theory
3. Lowpass Filters
4. Highpass Filters
5. Bandpass Filters


© 2000–2019 Joachim Weickert

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Motivation

Motivation

- ◆ In the previous lecture we have considered filters based on point operations. They ignore the neighbourhood structure of each pixel.
- ◆ Let us now study filters that can take into account the spatial context of the pixels.
- ◆ The simplest of these filters are linear and shift invariant. One can show that they can be described by convolutions.
- ◆ The behaviour of these filters can be nicely analysed in the Fourier domain: This turns convolutions into simple multiplications.
- ◆ Thus, linear shift invariant filters are elegant and transparent in theory and practise.

M I	
 A	
1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Linear System Theory

Linear Filter

- ◆ A filter L is called *linear*, if it satisfies the *superposition principle*

$$L(\alpha f + \beta g) = \alpha Lf + \beta Lg$$

for all (continuous or discrete) images f, g , and for all real numbers α, β .

Linear Shift Invariant (LSI) System

- ◆ A *shift (translation) invariant* filter acts identically at all locations.
- ◆ More formally, a shift-invariant filter L satisfies

$$LT_b f = T_b Lf$$

for all translations T_b with $(T_b f)(x) := f(x - b)$.

- ◆ A filter that is both linear and shift invariant is also called an *LSI system*.

Impulse Response of a Discrete LSI System

- ◆ The *impulse response (Impulsantwort)* of a discrete LSI filter L is the result of filtering a *discrete Dirac delta impulse*:

$$h = L\delta_0$$

where $\delta_0 = (\delta_{0,i})$ with the *Kronecker symbol*

$$\delta_{i,j} := \begin{cases} 1 & \text{for } i = j, \\ 0 & \text{else.} \end{cases}$$

- ◆ Since every discrete signal $f = (f_1, \dots, f_N)^\top$ can be represented as linear combination of N unit impulses $\delta_1, \dots, \delta_N$, linearity and shift invariance imply

$$Lf = L \sum_{i=1}^N f_i \delta_i = \sum_{i=1}^N f_i L\delta_i = \sum_{i=1}^N f_i L(T_i \delta_0) = \sum_{i=1}^N f_i T_i (L\delta_0).$$

- ◆ This shows: *Any LSI system L is fully characterised by its impulse response $L\delta_0$.*

Example: Stock Market Price Averaged over the Last 200 Days

$$u_i := \frac{1}{200} \sum_{k=0}^{199} f_{i-k}.$$

- ◆ This averaging can be represented as a discrete convolution (cf. Lecture 2):

$$u_i = \sum_{k=-\infty}^{\infty} f_{i-k} w_k = (\mathbf{f} * \mathbf{w})_i$$

with the convolution mask

$$w_k := \begin{cases} \frac{1}{200} & \text{for } k \in \{0, \dots, 199\}, \\ 0 & \text{else.} \end{cases}$$

- ◆ Such a convolution filter is linear and shift invariant (Assignment H1, Problem 3).
- ◆ Its impulse response $\mathbf{h} = \mathbf{L}\delta_0 = \delta_0 * \mathbf{w}$ is given by the convolution mask:

$$h_i = \sum_{k=-\infty}^{\infty} \delta_{0,i-k} w_k = w_i \quad \forall i.$$

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Repetition from Lecture 2: Convolution

- ◆ discrete convolution in 1-D:

$$(\mathbf{f} * \mathbf{w})_i := \sum_{k=-\infty}^{\infty} f_{i-k} w_k$$

- ◆ discrete convolution in 2-D:

$$(\mathbf{f} * \mathbf{w})_{i,j} := \sum_{k=-\infty}^{\infty} \sum_{\ell=-\infty}^{\infty} f_{i-k, j-\ell} w_{k,\ell}$$

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Linear System Theory (5)



- ◆ continuous convolution in 1-D:

$$(f * w)(x) := \int_{-\infty}^{\infty} f(x-x') w(x') dx'$$

- ◆ continuous convolution in 2-D:

$$(f * w)(x, y) := \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x-x', y-y') w(x', y') dx' dy'$$

Signals with finite extension can be mirrored and extended periodically.

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Linear System Theory (6)



Important Properties of the Convolution

(cf. Homework H1, Problem 3, and Lecture 4)

- ◆ **Linearity:**

$$(\alpha f + \beta g) * w = \alpha (f * w) + \beta (g * w)$$

for all signals/images f, g and all real numbers α, β .

- ◆ **Shift Invariance:**

$$T_b(f * w) = (T_b f) * w$$

for all translations T_b .

- ◆ **Commutativity:**

$$f * w = w * f.$$

Function and convolution kernel play an equal role.

- ◆ **Associativity:**

$$(f * v) * w = f * (v * w).$$

Successive convolution with kernels v and w comes down to a single convolution with the kernel $v * w$.

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Linear System Theory (7)



◆ Distributivity:

$$(f + g) * w = f * w + g * w.$$

◆ Differentiation:

$$(f * w)' = f' * w = f * w'.$$

Thus, either the signal or the kernel is differentiated.

◆ Differentiability:

Convolving with a smooth kernel makes a signal smoother:

If $f \in C^0(\mathbb{R})$ and $w \in C^n(\mathbb{R})$, then $(f * w) \in C^n(\mathbb{R})$.

◆ Convolution Theorem of the Fourier Transform:

$$\mathcal{F}[f * w] = \mathcal{F}[f] \cdot \mathcal{F}[w].$$

This allows an efficient convolution if the kernels have a large support region.

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Linear System Theory (8)



Importance of Convolutions in Linear System Theory

- ◆ Any convolution is linear and shift invariant, i.e. it creates an LSI system.
- ◆ More importantly, it can be shown that even the reverse is true:
An LSI system always performs a convolution !
- ◆ The convolution mask is given by the impulse response.

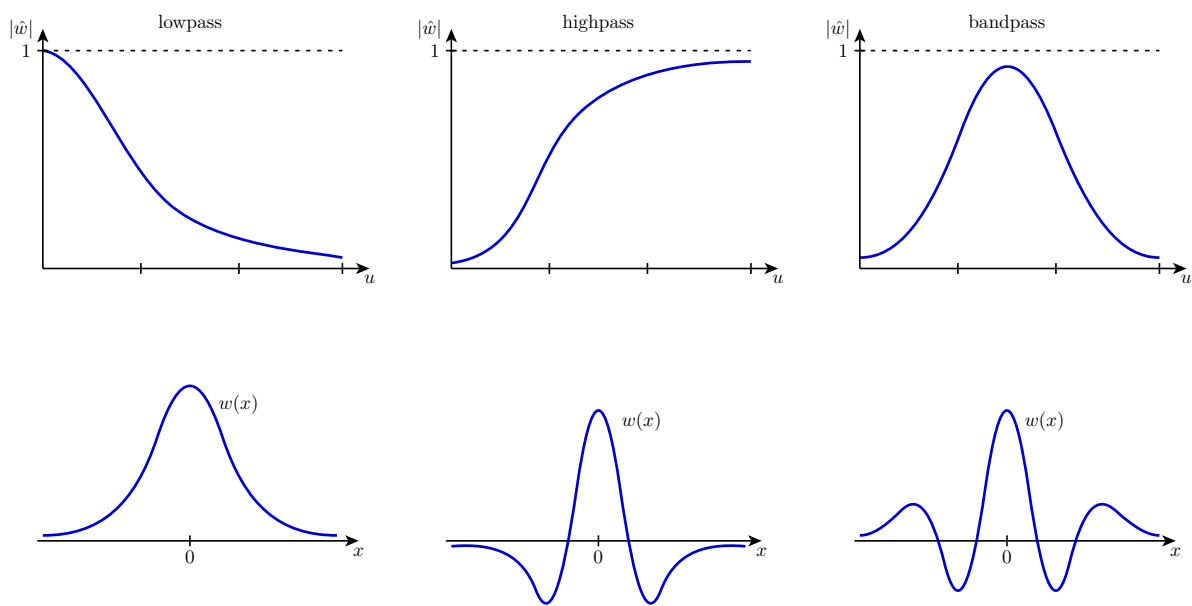
Importance of the Fourier Transform in Linear System Theory

- ◆ We have seen that LSI systems are fully characterised by convolutions.
- ◆ Convolutions in the spatial domain become multiplications in the Fourier domain. Thus, Fourier analysis is perfectly suited for LSI filters.
- ◆ For large convolution kernels, it is more efficient to perform the computation in the Fourier domain.
- ◆ LSI filters are often studied in the Fourier domain, in order to understand their frequency behaviour.
- ◆ Often one even starts designing LSI filters in the Fourier domain. Afterwards one transforms them back to the spatial domain.

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Basic Types of LSI Filters



Top: Basic types of convolution kernels in the Fourier domain. **Bottom:** Corresponding kernels in the spatial domain. Author: T. Schneevoigt.

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Often one uses the following taxonomy to characterise LSI filters:

- ◆ **Lowpass filters:** Low frequencies are less attenuated than high ones.
- ◆ **Highpass filters:** High frequencies are less attenuated than low ones.
- ◆ **Bandpass filters:** A specific frequency band is hardly attenuated.

Let us now study these three LSI filter types in more detail.

Lowpass Filters

Goals

- ◆ smooth an image by eliminating noise and unimportant small-scale details
- ◆ design in the spatial domain: convolution with a weighted averaging mask
- ◆ design in the Fourier domain: attenuate high frequencies

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Design in the Spatial Domain: Box Filters

- ◆ uses convolution mask of size $(2m+1) \times (2m+1)$ with weights $\frac{1}{(2m+1)^2}$
- ◆ can also be implemented efficiently for large masks in the spatial domain:
 - filter is separable
 - Shifting the 1-D mask by one pixel to the right removes one grey value at the left end, and it adds one at the right end:

$$\begin{aligned}
 u_{i+1} &= \frac{1}{2m+1} \sum_{k=-m}^m f_{i+1-k} \\
 &= \frac{1}{2m+1} \left(f_{i+m+1} - f_{i-m} + \sum_{k=-m}^m f_{i-k} \right) \\
 &= \frac{1}{2m+1} \left(f_{i+m+1} - f_{i-m} \right) + u_i
 \end{aligned}$$

- total complexity is linear and independent (!) of the mask size:
1 addition, 1 subtraction, 1 multiplication per pixel (in 1-D)

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Lowpass Filters (3)



- ◆ Often results with a single box filter do not look too convincing:
 - not rotationally invariant:
 - prefers horizontal and vertical structures
 - not satisfactory in the frequency domain:
 - continuous FT of a box function is a sinc function (Lecture 4)
 - nonmonotone behaviour: Fourier spectrum has multiple extrema
 - attenuation of high frequencies only with $1/|u|$
- ◆ However, we will see later that iterated box filtering can be useful.

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Lowpass Filters (4)



Optimality in the Fourier Domain: The Ideal Lowpass

- ◆ All frequency components (u, v) with $u^2 + v^2 > T^2$ are set to 0.
- ◆ not satisfactory in the spatial domain:
 - sinc-like rotation invariant convolution kernel in the spatial domain
 - creates visible ringing artifacts

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Lowpass Filters (5)

Gaussian Convolution Kernels

- ◆ m -dimensional Gaussian:


$$K_{\sigma}(\mathbf{x}) := \frac{1}{(2\pi\sigma^2)^{m/2}} \exp\left(-\frac{|\mathbf{x}|^2}{2\sigma^2}\right).$$

The “width” σ is called *standard deviation*, σ^2 is the *variance*.

- ◆ creates Gaussian with reciprocal variance in Fourier domain (Lecture 4)
- ◆ good compromise: one maximum in both spatial and frequency domain
- ◆ the only convolution kernel that is both separable and rotationally invariant
- ◆ Iterated Gaussian convolution creates a new Gaussian where the variances sum up.
- ◆ Gaussian convolution can be implemented efficiently in numerous ways.

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Lowpass Filters (6)



Top left: Noisy original image. **Top right:** Filtering with a 11×11 box filter creates horizontal and vertical artifacts. **Bottom left:** The ideal lowpass with $T^2 = 500$ suffers from ringing artifacts. **Bottom right:** Smoothing with a Gaussian with $\sigma = 3$ gives better results. Author: J. Weickert.

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Lowpass Filters (7)



Approximation Possibility 1: Sampling in the Spatial Domain

- ◆ exploit separability and symmetry in order to achieve high efficiency
- ◆ restrict sampling to interval $[-k\sigma, k\sigma]$ (high accuracy for $k \geq 3$)
- ◆ renormalise sum of coefficients to 1
- ◆ Advantage: simple and flexible (σ can be tuned continuously)
- ◆ Disadvantage: computational complexity increases with σ
- ◆ good for small values of σ

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Lowpass Filters (8)



Approximation Possibility 2: Multiplication in the Fourier Domain

- ◆ cf. Lecture 5:
 - use FFT to transform the image into the Fourier domain
 - multiply with the Fourier transform of the Gaussian (a Gaussian with inverse variance)
 - use FFT for backtransformation
- ◆ Advantages:
 - almost linear complexity: $\mathcal{O}(N^2 \log N)$ for an $N \times N$ image
 - computational complexity does not increase with σ
- ◆ Disadvantages:
 - wraparound errors (unless image is mirrored)
 - standard FFT requires image sizes of powers of 2
- ◆ good for large values of σ

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Lowpass Filters (9)



Approximation Possibility 3: Binomial Kernels

Binomial kernels and the variance of the approximated Gaussians.

Normalisation	Filter Coefficients	Variance σ^2
1	1	0
1/2	1 1	1/4
1/4	1 2 1	1/2
1/8	1 3 3 1	3/4
1/16	1 4 6 4 1	1
1/32	1 5 10 10 5 1	5/4
1/64	1 6 15 20 15 6 1	3/2
1/128	1 7 21 35 35 21 7 1	7/4
1/256	1 8 28 56 70 56 28 8 1	2

- ◆ Binomial kernels approximate Gaussians.
- ◆ Separability and symmetry can be exploited.
- ◆ Iterated binomial kernels create binomial kernels.

Lowpass Filters (10)



- ◆ Advantage:
 - even possible in integer arithmetics:
division by powers of 2 comes down to bit shifts
- ◆ Disadvantages:
 - computational complexity increases with σ
 - σ cannot be tuned continuously

Approximation Possibility 4: Iterated Box Filters




Illustration of the central limit theorem of statistics. **Left:** Iterated box filters approximate a Gaussian. The graph shows the box filter (black) and the results after 1 (red), 2 (green), and 5 (blue) iterations. **Right:** Iterating a more complicated filter also approximates a Gaussian, but the convergence is slower. The graph shows a parabola-shaped filter (black) and the results after 1 (red), 2 (green), and 10 (blue) iterations. Author: T. Schneevoigt.

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Lowpass Filters (12)

- ◆ central limit theorem of statistics:
iterated averaging kernels (symmetric, normalised, nonnegative) converge to Gaussians
- ◆ Iterating a box filter three times is already a reasonable approximation.
- ◆ It can be shown that n iterations of a box filter $(b_i)_{i \in \mathbb{Z}}$ of length $2\ell + 1$,

$$b_i = \begin{cases} \frac{1}{2\ell+1} & \text{for } -\ell \leq i \leq \ell, \\ 0 & \text{else,} \end{cases}$$

approximate a Gaussian with variance

$$\sigma^2 = n \cdot \frac{\ell^2 + \ell}{3}.$$

- ◆ Advantage: linear complexity, independent of σ
- ◆ Disadvantage: σ cannot be tuned continuously

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Highpass Filters

Goals

- ◆ remove low-frequent background perturbations
- ◆ perhaps even sharpen blurry image structures by enhancing high frequencies

Remarks

- ◆ An important class of highpass filters consists of derivative filters. They are useful for detecting edges (next lecture).
- ◆ While lowpass filters act stabilising, highpass filters may act destabilising, if they enhance high frequencies.

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Design in the Spatial Domain

- ◆ Example: highpass filter as difference between identity and a lowpass filter
- ◆ Using e.g. a 3×3 box filter as lowpass filter creates the highpass stencil

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 1 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} - \begin{array}{|c|c|c|} \hline \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \hline \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \hline \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline -\frac{1}{9} & -\frac{1}{9} & -\frac{1}{9} \\ \hline -\frac{1}{9} & \frac{8}{9} & -\frac{1}{9} \\ \hline -\frac{1}{9} & -\frac{1}{9} & -\frac{1}{9} \\ \hline \end{array}$$

where the *stencil notation* depicts the weights in the pixels

$i-1, j+1$	$i, j+1$	$i+1, j+1$
$i-1, j$	i, j	$i+1, j$
$i-1, j-1$	$i, j-1$	$i+1, j-1$


1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32



Highpass Filters (3)

M	I
	A
1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32


- ◆ Displaying the filtered image often requires an affine rescaling of the grey values (cf. Lecture 10):
 - For many of these filters, the average grey value becomes 0 (e.g. if the lowpass filter preserves the average grey value).
 - Thus, negative values are common.

Highpass Filters (4)

M	I
	A
1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32




Left: Vessel structure of the background of the eye. **Right:** Elimination of low-frequency background structures by subtracting a Gaussian-smoothed version from the original image. The greyscale range $[-94, 94]$ has been rescaled to $[0, 255]$ by an affine rescaling. Author: J. Weickert.

M	I
	A
1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

Bandpass Filters

- ◆ useful for extracting interesting image structures on certain scales
- ◆ Example: assessing the cloudiness of fabrics (Lecture 6)
- ◆ can be created by subtracting two lowpass filters
- ◆ If the lowpass filters are Gaussians, the resulting bandpass is called *DoG (difference of Gaussians)*.

M	I
	A
1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32



(a) Top left: Fabric, 257×257 pixels. (b) Top right: After lowpass filtering with a Gaussian with $\sigma = 10$. (c) Bottom left: Lowpass filtering with $\sigma = 15$. (d) Bottom right: Subtracting (b) and (c) gives a bandpass filter that visualises cloudiness on a certain scale. The greyscale range has been affinely rescaled from $[-13, 13]$ to $[0, 255]$. Author: J. Weickert.

Summary

- ◆ Linear shift invariant (LSI) filters are fully characterised by their impulse response.
- ◆ can always be represented as convolutions
- ◆ impulse response: given by convolution mask
- ◆ The Fourier transform is very important for designing LSI filters.
- ◆ Lowpass filters are useful for smoothing data:
 - most important example: Gaussian convolution
 - Gaussian convolution can be implemented in many ways, e.g. in the spatial domain, in the Fourier domain, via binomial filters, via iterated box filters
- ◆ Highpass filters eliminate low-frequent perturbations and/or sharpen image structures.
- ◆ Bandpass filters are mainly used for extracting features at certain scales.

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32

References

- ◆ K. R. Castleman: *Digital Image Processing*. Prentice Hall, Upper Saddle River, 1996.
(a text book that focuses on LSI filters)
- ◆ R. C. Gonzalez, R. E. Woods: *Digital Image Processing*. Pearson, Upper Saddle River, Global Edition, 2017.
(see in particular Sections 3.5–3.7 and Chapter 4)
- ◆ A. V. Oppenheim, R. W. Schaffer: *Discrete-Time Signal Processing*. Prentice Hall, Englewood Cliffs, Third Edition, 2013.
(Signal processing books usually provide an exhaustive treatment of LSI filters. This volume is one of the classical text books on signal processing.)
- ◆ W. M. Wells: Efficient synthesis of Gaussian filters by cascaded uniform filters. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, Vol. 8, pp. 234–239, 1986.
(Gaussian approximation by iterated box filters)
- ◆ P. Getreuer: A survey of Gaussian convolution algorithms. *Image Processing On Line*, Vol. 3, pp. 286–310, 2013.
(<https://www.ipol.im/pub/art/2013/87/>)
(compares many algorithms for Gaussian convolution)

1	2
3	4
5	6
7	8
9	10
11	12
13	14
15	16
17	18
19	20
21	22
23	24
25	26
27	28
29	30
31	32